

Investigating the Proactor design pattern

Martin Matusiak

November 29, 2007



Outline

What it's all about

I/O strategies

The case for Proactor

But didn't we have Reactor for that?

But is it any good?

Proactors in the wild?



What to expect from this talk

- ▶ My impression: Proactor is a complicated pattern.
- ▶ I do not follow the textbook, instead I give you my thoughts/conclusions based on some extra background reading.
- ▶ Hoping to fill in the blanks you may have had while reading about it.
- ▶ Proactor seems very similar to Reactor, I go in depth on this.
- ▶ Questions underway? Feel free to interrupt.



Based on these sources

1. *Pattern-Oriented Software Architecture : Patterns for Concurrent and Networked Objects* [our textbook]
2000, Schmidt et al.
2. *Proactor: An Object Behavioral Pattern for Demultiplexing and Dispatching Handlers for Asynchronous Events*
1997, Schmidt et al.
3. *Applying the Proactor Pattern to High-Performance Web Servers*
1998, Schmidt et al.
4. *ProActor vs Reactor: Comparing Two High-Performance I/O Design Patterns*
2005?, Libman & Gilbourn

Characteristics

- ▶ Quite old and possibly outdated sources.
- ▶ Seemingly little interest and/or knowledge of Proactor/Reactor patterns outside Schmidt and friends.



What is it good for?

- ▶ I/O! (no other use case given in sources)



What is it good for?

- ▶ I/O! (no other use case given in sources)
- ▶ *More specifically*: how to process I/O events from multiple sources concurrently.



What is it good for?

- ▶ I/O! (no other use case given in sources)
- ▶ *More specifically*: how to process I/O events from multiple sources concurrently.



- ▶ *(++More) specifically*: how to decouple I/O concurrency from process concurrency (n:m relation)



Outline

What it's all about

I/O strategies

The case for Proactor

But didn't we have Reactor for that?

But is it any good?

Proactors in the wild?



Synchronous, blocking I/O

```
lines = read("info.txt");  
//  thread blocks until I/O call completes  
//  moved to I/O wait list, not runnable  
  
//  * wait until call completes *  
  
//  thread status changed back to runnable  
print(lines);
```



Synchronous, non-blocking I/O

```
myhandle = from_filename("info.txt");

while (!lines) {
    result = select(myhandle, timeout);
    // thread switched to I/O bound

    // * wait until call completes *

    // thread switched back to cpu bound

    if (result) {
        lines = read(handle);
        // syscall has confirmed handle is ready, I/O call
        // is synchronous and non-blocking
    }
}
print(lines);
```



Asynchronous I/O

```
int main {  
    read("info.txt");  
    // call is asynchronous, thread still cpu bound  
    // I/O operation executes in kernel thread,  
    // not application thread  
}  
  
void handle_read(string lines) {  
    print(lines);  
}
```

Missing some more OS-specific plumbing, but this is the basic idea.
Note: call site differs from return site.



Outline

What it's all about

I/O strategies

The case for Proactor

But didn't we have Reactor for that?

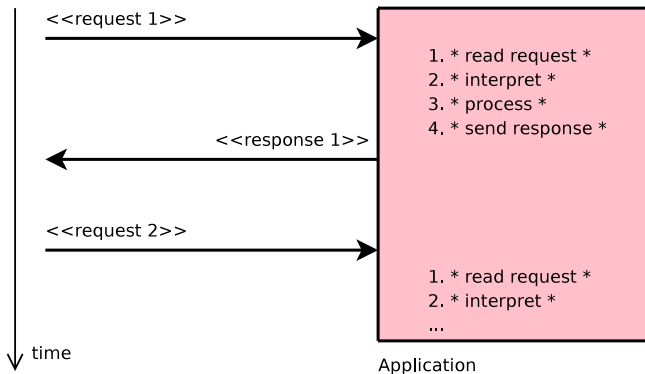
But is it any good?

Proactors in the wild?




Handling I/O in a web server

Single source, sequential I/O, **“one at a time”**
→ single thread



Handling I/O in a web server

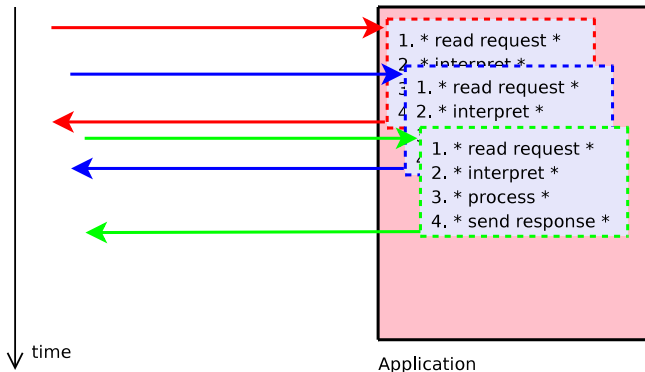
Multiple source, concurrent I/O, “**many at a time**”



Handling I/O in a web server

Multiple source, concurrent I/O, “**many at a time**”

→ thread per request, *multi-threaded synchronous I/O*

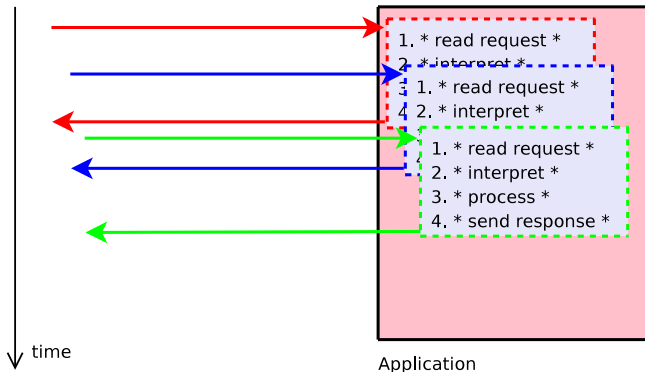


Handling I/O in a web server

Multiple source, concurrent I/O, “**many at a time**”

→ thread per request, *multi-threaded synchronous I/O*

✓ requests are independent (threads have no shared state)



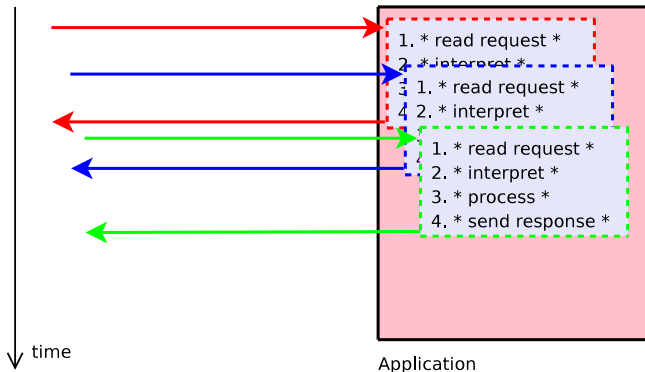
Handling I/O in a web server

Multiple source, concurrent I/O, “**many at a time**”

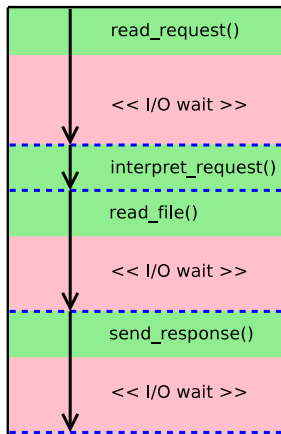
→ thread per request, *multi-threaded synchronous I/O*

✓ requests are independent (threads have no shared state)

☠ context switching



Multi-threaded synchronous I/O

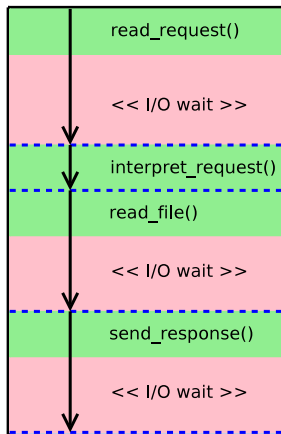


The life of a thread

- ▶ thread lifetime = request processing duration



Multi-threaded synchronous I/O

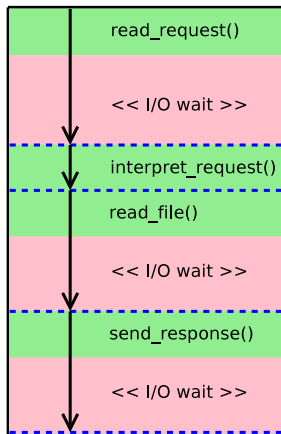


The life of a thread

- ▶ thread lifetime = request processing duration
- ▶ spends most of the time waiting



Multi-threaded synchronous I/O

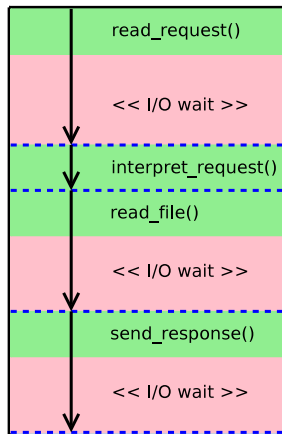


The life of a thread

- ▶ thread lifetime = request processing duration
- ▶ spends most of the time waiting
- ▶ *has to be in memory*



Multi-threaded synchronous I/O



The life of a thread

- ▶ thread lifetime = request processing duration
- ▶ spends most of the time waiting
- ▶ *has to be in memory*
- ▶ *has to be runnable (ie. subject to context switches)**



Outline

What it's all about

I/O strategies

The case for Proactor

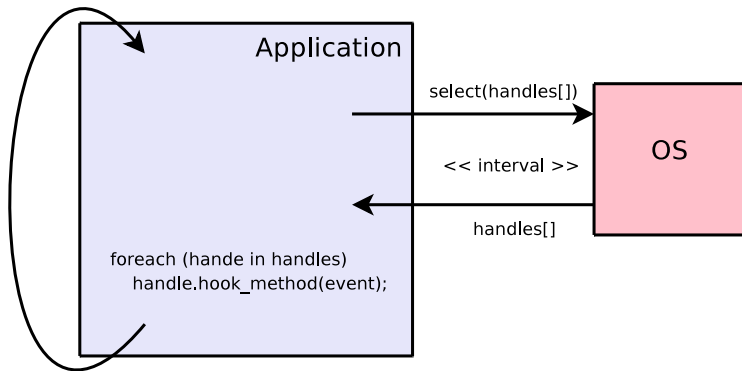
But didn't we have Reactor for that?

But is it any good?

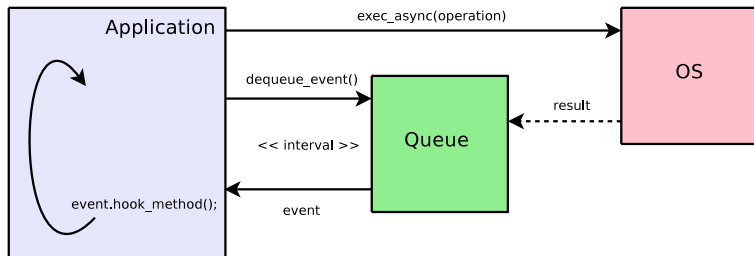
Proactors in the wild?



The Reactor pattern



The Proactor pattern




So what is the difference exactly?

- ▶ The patterns are very similar.
 - A single thread of execution.
 - Multiple I/O operations ongoing simultaneously.
 - Execution of event handlers is sequential.
- ▶ They differ in minor details.
 - Proactor *launches* (or even generates) I/O events and waits for callbacks, ie. **proactive stance**.
Reactor only *anticipates* [external] events, ie. **reactive stance**.
 - Proactor delegates I/O operations to the OS – **asynchronous**
Reactor runs I/O operations in its own thread – **synchronous**




Reactor: synchronous I/O



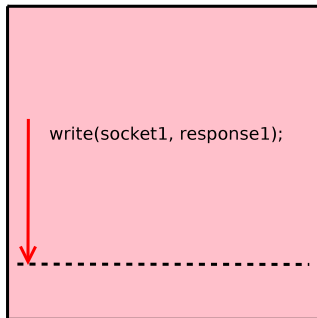
Application



Proactor: asynchronous I/O



Application



Operating system



Comparison revisited

- ▶ Pattern wise they are almost the same.
 - You could even say they are compositions of each other.
TProactor (Terabit Solutions) is a proactor emulated on top of a reactor. (no performance degradation)



Comparison revisited

- ▶ Pattern wise they are almost the same.
 - You could even say they are compositions of each other.
TProactor (Terabit Solutions) is a proactor emulated on top of a reactor. (no performance degradation)
- ▶ The difference is in how they use OS services for I/O.
 - The distinction is therefore in use of OS, not in design.
 - Proactor is just a way of taking advantage of async I/O support in the OS.



Comparison revisited

- ▶ Pattern wise they are almost the same.
 - You could even say they are compositions of each other.
TProactor (Terabit Solutions) is a proactor emulated on top of a reactor. (no performance degradation)
- ▶ The difference is in how they use OS services for I/O.
 - The distinction is therefore in use of OS, not in design.
 - Proactor is just a way of taking advantage of async I/O support in the OS.
- ▶ Is asynchronous I/O faster than synchronous I/O? If so, why?
 - Kernel design territory, not application/middleware design.



Outline

What it's all about

I/O strategies

The case for Proactor

But didn't we have Reactor for that?

But is it any good?

Proactors in the wild?




A web server based benchmark

- ▶ Source: *Applying the Proactor Pattern to High-Performance Web Servers*
- ▶ Experiment setup:
 - JAWS web server with pluggable concurrency strategies:
 - thread per request
 - thread pool
 - proactor pattern
 - Method: generate `GET` requests for documents of a given size, steadily increase number of concurrent clients.
 - Measurements:
 - Throughput (data per unit of time)
 - Latency (time elapsed from request to response)



Benchmark: 5kb files

Throughput: higher is better



Latency: lower is better




Figure 7: Experiment Results from 5K File

Most common size of documents on a typical web server.



Benchmark: 50kb files





Figure 8: Experiment Results from 50K File

Second most common size of documents on a typical web server.



Benchmark: 500kb files





Figure 9: Experiment Results from 500K File

Break point reached, proactor based profile (TransmitFile) has highest throughput and lowest latency.



Benchmark: 5mb files





Figure 10: Experiment Results from 5M File



Outline

What it's all about

I/O strategies

The case for Proactor

But didn't we have Reactor for that?

But is it any good?

Proactors in the wild?



Synchronous non-blocking I/O

- ▶ `select` **nix, BSD, Windows*
- ▶ `poll` **nix, BSD*
- ▶ `epoll` (maintains state, more efficient) *Linux*
- ▶ `KQueue` (similar to `epoll`) *BSD*
- ▶ `/dev/poll` *Solaris*
- ▶ `libevent` (encapsulates OS-specific interfaces) *Linux, BSD, Solaris, Windows*



Synchronous non-blocking I/O

- ▶ `select` **nix, BSD, Windows*
- ▶ `poll` **nix, BSD*
- ▶ `epoll` (maintains state, more efficient) *Linux*
- ▶ `KQueue` (similar to `epoll`) *BSD*
- ▶ `/dev/poll` *Solaris*
- ▶ `libevent` (encapsulates OS-specific interfaces) *Linux, BSD, Solaris, Windows*

Asynchronous I/O

- ▶ Completion ports, overlapped I/O *Windows NT* (current?)
- ▶ POSIX AIO *Linux, BSD, others?*
- ▶ `io_submit()`, `io_getevents()` *Linux*



Twisted framework for Python

- ▶ *“Twisted is an event-driven networking engine written in Python”*
- ▶ Borrows heavily from the ACE framework (Schmidt et al.)



Twisted framework for Python

- ▶ *“Twisted is an event-driven networking engine written in Python”*
- ▶ Borrows heavily from the ACE framework (Schmidt et al.)
- ▶ *Reactor* pattern used as one of the basic building blocks.



Twisted framework for Python

- ▶ *“Twisted is an event-driven networking engine written in Python”*
- ▶ Borrows heavily from the ACE framework (Schmidt et al.)
- ▶ *Reactor* pattern used as one of the basic building blocks.
- ▶ Originally created as a reusable network library for game development.



Twisted framework for Python

- ▶ *“Twisted is an event-driven networking engine written in Python”*
- ▶ Borrows heavily from the ACE framework (Schmidt et al.)
- ▶ *Reactor* pattern used as one of the basic building blocks.
- ▶ Originally created as a reusable network library for game development.
- ▶ *Twisted Core* is an async event loop geared at fast development of various async client and server applications.



Twisted framework for Python

- ▶ “*Twisted is an event-driven networking engine written in Python*”
- ▶ Borrows heavily from the ACE framework (Schmidt et al.)
- ▶ *Reactor* pattern used as one of the basic building blocks.
- ▶ Originally created as a reusable network library for game development.
- ▶ *Twisted Core* is an async event loop geared at fast development of various async client and server applications.

No uses of *Proactor* pattern found. ☹



- ▶ Is the pattern clear to you? Is it straightforward or is it subtle?
- ▶ Does the complexity of the Proactor pattern intimidate you?
- ▶ Do you believe Reactor and Proactor differ significantly beyond the evidence presented?

